

# NBA Basketball and Rookies – Regularised Regression

Valérie BLANCH

2614867B

## 1. INTRODUCTION

The dataset analysed here contains statistics on 600 basketball players during their rookie year at the NBA: the question of interest that justifies this project is whether it is possible to predict the career length of the players from their first year's performance, and which variable, or combination of variables will be the most useful during this analysis.

The dataset is structured as follows:

```
> str(basket)
'data.frame': 600 obs. of 23 variables:
 $ Name      : chr  "Mustafa Shakur" "Alec Burks" "Charles Thomas" "J.R. Henderson"
 ...
 $ Year_drafted: int  2010 2011 1991 1998 1991 1987 1989 2000 2009 1989 ...
 $ GP         : int  22 59 36 30 51 40 50 50 25 55 ...
 $ MIN        : num  7.2 15.9 4.3 11 20.7 9.5 9.8 18.7 5.2 18 ...
 $ PTS        : num  2.3 7.2 1.3 3.2 8.4 3.8 5.2 6.3 2.3 6.1 ...
 $ FG_made    : num  1 2.6 0.5 1.2 3.2 1.6 1.9 2.3 0.8 2.5 ...
 $ FGA        : num  2.7 6.1 1.4 3.2 6.4 3.7 4.1 6.2 2.1 5.2 ...
 $ FG_percent : num  35.6 42.9 35.3 36.5 49.7 42.2 47.3 37 39.6 47.9 ...
 $ TP_made    : num  0 0.3 0.1 0.1 0 0.1 0.1 0.6 0.2 0 ...
 $ TPA        : num  0.5 0.8 0.5 0.2 0 0.3 0.3 1.7 0.6 0 ...
 $ TP_percent : num  10 33.3 11.8 40 0 20 35.7 33.3 31.3 0 ...
 $ FT_made    : num  0.4 1.8 0.3 0.8 2 0.7 1.2 1.1 0.4 1.2 ...
 $ FTA        : num  0.7 2.4 0.4 1.5 2.9 0.7 1.8 1.4 0.5 1.7 ...
 $ FT_percent : num  53.3 72.7 66.7 55.6 66.7 89.7 67 77.8 84.6 66.7 ...
 $ OREB       : num  0.3 1 0.2 0.7 2.3 0.1 0.6 0.4 0.3 1.5 ...
 $ DREB       : num  0.7 1.3 0.4 0.9 3.1 0.6 0.9 1.2 0.4 1.9 ...
 $ REB        : num  1 2.2 0.6 1.6 5.4 0.7 1.5 1.6 0.7 3.4 ...
 $ AST        : num  1.1 0.9 0.6 0.7 0.8 1.5 0.9 2.9 0.5 1 ...
 $ STL        : num  0.2 0.5 0.1 0.3 0.5 0.1 0.6 1 0.3 0.4 ...
 $ BLK        : num  0.1 0.1 0 0.1 0.5 0.1 0 0 0.1 0.2 ...
 $ TOV        : num  0.8 0.9 0.5 0.6 1.5 0.6 0.8 1.1 0.6 1.1 ...
 $ Yrs        : int  2 5 1 1 5 2 3 1 4 3 ...
 $ Target     : int  0 0 0 0 0 0 0 0 0 0 ...
```

- 600 rows corresponding to the 600 players.
- The last column of the data frame is the target variable, which is a binary outcome: "0" if the player's career length was less than or equal to 5 years, "1" otherwise.
- 22 variables containing various statistics for each player, which will be the predictors of the models:
  - o **Categorical variables:** *Name* (the name of the player), *Year\_drafted* (the year the player was drafted).
  - o **Numeric variables:** *GP* (Games Played - out of 82), *PTS* (Points per game), *FG\_made* (Field goals made - per game), *FGA* (Field goal attempts - per game), *TP\_made* (Three points made - per game), *TPA* (Three point attempts - per game), *FT\_made* (Free throws made - per game), *FTA* (Free throws attempts - per game), *OREB* (Offensive rebounds - per game), *DREB* (Defensive rebounds - per game), *REB* (Total rebounds - per game), *AST* (Assists - per game), *STL* (Steals - per game), *BLK* (Blocks - per game), *TOV* (Turnovers - per game), *MIN* (Minutes per game - out of 48), *FG\_percent* (Field goal percentage), *TP\_percent* (Three point percentage), *FT\_percent* (Free throws percentage), *Yrs* (Career length - in years).

## 2. EXPLORATORY ANALYSIS

*Pre-processing:*

- Both the columns *Name* and *Yrs* have been dropped: the former because this variable is not relevant for this regression analysis, and the latter because it duplicates the target variable and thus produces perfect separation.
- The variables *Target* and *Year\_drafted* have been converted to factors.
- The dataset has been tested for missing values, but there were none.

```
> anyNA(basket)
[1] FALSE
```

- The dataset has been separated into a training set (70%) and a test set (30%) using the `createDataPartition()` function from the *caret* package: that way both categories of the target variable are evenly distributed in each set. The training set contains 21 variables and 420 observations, while the test set contains 21 variables for 180 observations.


```
> dim(train.set)
[1] 420 21

> dim(test.set)
[1] 180 21
```

- The values have been scaled since they were recorded in different units.

*EDA of the training set:*


As mentioned above, the two categories of the target variable are evenly distributed in the training set.



210 players have a career less than or equal to 5 years and 210 have a career greater than 5 years, as shown in this proportion table:

```
> table(train.set$Target)
 0   1
210 210
```

After calculating the correlation between the numeric variables, a plot has been drawn and shows obvious evidence of multicollinearity.



The multicollinearity between the variables of the training set justifies the use of regularised logistic models.

Three different techniques have been used:

- Lasso regression, which sets the irrelevant or correlated variables to zero, thus provides a strict variable selection and highlights which are the most influential.
- Ridge regression, which weights down irrelevant or correlated variables, thus keeps all of the variables and avoids any potential loss of information.
- Elastic net regression, which incorporates both algorithms.

### 3. RESULTS

All of the models in this study have been fitted using the *glmnetUtils* package, which provides the same functions as the *glmnet* one, plus the `cva.glmnet()` function which calculates the optimal alpha parameter for elastic net models.

*Lasso Regression:*

First, a lasso model has been computed using the following function and arguments:

```
> cv.lasso <- cv.glmnet(x, y, alpha=1, type.measure='class',
+ family='binomial')
```

The `cv.glmnet()` function computes a 10-folds cross-validation by default. The `alpha` parameter has been set to 1 for lasso regression, the `family` option has been set to 'binomial' for the algorithm to recognize the target as binary. The `type.measure` argument is set to 'class': the function will find the best lambda (shrinkage parameter) according to the lowest misclassification error.

```
Measure: Misclassification Error

      Lambda Index Measure      SE Nonzero
min 0.03587   19  0.3095 0.0207      8
1se 0.13196    5  0.3238 0.0249      3
```

The function returns two lambda parameters: *lambda.min*, which is the exact value of lambda, and *lambda.1se*, which simplifies the model further within 1 standard error. Thus, the model using *lambda.min* has a lower misclassification error rate of 0.31 and 8 non-zero variables, and the one using *lambda.1se* has a higher misclassification error rate but has kept only 3 variables, *GP*, *MIN* and *REB*.

```
GP      0.1283905650
MIN      0.0140638773
PTS      .
FG_made  .
FGA      .
FG_percent .
TP_made  .
TPA      .
TP_percent .
FT_made  .
FTA      .
FT_percent .
OREB     .
DREB     .
REB      0.1715358220
```

Figure 1 : `coef(cv.lasso)` output (cropped)

### Predictions:

```
> #Lasso Model with lambda.min
> confusionMatrix(as.factor(pred.lasso.min))
Confusion Matrix and Statistics
```

```
      Reference
Prediction 0  1
0  68  28
1  22  62
```

```
Accuracy : 0.7222
95% CI : (0.6507, 0.7863)
No Information Rate : 0.5
P-Value [Acc > NIR] : 1.067e-09
```

```
Kappa : 0.4444
```

```
McNemar's Test P-Value : 0.4795
```

```
Sensitivity : 0.6889
Specificity : 0.7556
Pos Pred Value : 0.7381
Neg Pred Value : 0.7083
Prevalence : 0.5000
Detection Rate : 0.3444
Detection Prevalence : 0.4667
Balanced Accuracy : 0.7222
```

```
'Positive' Class : 1
```

```
> #Lasso Model with lambda.1se
> confusionMatrix(as.factor(pred.lasso.1se))
Confusion Matrix and Statistics
```

```
      Reference
Prediction 0  1
0  69  32
1  21  58
```

```
Accuracy : 0.7056
95% CI : (0.6332, 0.771)
No Information Rate : 0.5
P-Value [Acc > NIR] : 1.716e-08
```

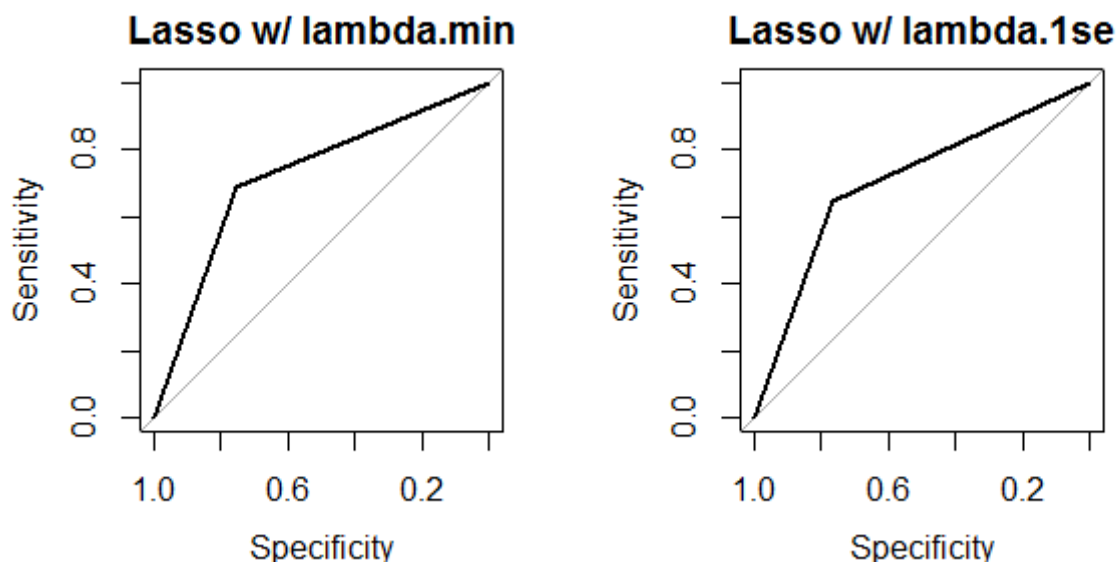
```
Kappa : 0.4111
```

```
McNemar's Test P-Value : 0.1696
```

```
Sensitivity : 0.6444
Specificity : 0.7667
Pos Pred Value : 0.7342
Neg Pred Value : 0.6832
Prevalence : 0.5000
Detection Rate : 0.3222
Detection Prevalence : 0.4389
Balanced Accuracy : 0.7056
```

```
'Positive' Class : 1
```

Predictions have been computed using the `predict()` function, the test set, and both *lambda.min* and *lambda.1se*. The outputs of the `confusionMatrix()` function (*caret* package) are shown above. The overall accuracy of the lasso model with shrinkage parameter *lambda.min* is greater than the one with *lambda.1se* (0.72 vs 0.71), although the difference is not significant. Both models have a specificity greater than the sensitivity, meaning they both have trouble detecting positive outcomes (false negatives or type II errors). This slight imbalance between type I and type II errors can be seen with ROC curves (*pROC* package):



Therefore, considering the outputs above, the lasso model with *lambda.min* as a shrinkage parameter is the most performant one, although *lambda.1se* could be used as well, since its metrics are not much different, and its model's simplicity might be preferable for inference.

### Ridge Regression:

Two ridge regression models have been fitted, using the same function and same arguments, apart from *alpha*, which has been set to 0. The `cv.glmnet()` function has calculated the following lambda parameters:

Measure: Misclassification Error					
	Lambda	Index	Measure	SE	Nonzero
min	2.201	49	0.3167	0.02536	51
1se	14.150	29	0.3310	0.02805	51

The misclassification errors are similar to the lasso models', slightly above 0.30. No variable has been set to zero, as expected from a ridge model.

The same functions and arguments have been used to compute the following predictions and metrics:

```
> #Ridge model w/ lambda.min
```

```
> confusionMatrix(as.factor(pred.ridge.min))
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	72	27
1	18	63

Accuracy : 0.75  
 95% CI : (0.6801, 0.8114)  
 No Information Rate : 0.5  
 P-Value [Acc > NIR] : 6.046e-12

Kappa : 0.5

Mcnemar's Test P-value : 0.233

Sensitivity : 0.7000  
 Specificity : 0.8000  
 Pos Pred Value : 0.7778  
 Neg Pred Value : 0.7273  
 Prevalence : 0.5000  
 Detection Rate : 0.3500  
 Detection Prevalence : 0.4500  
 Balanced Accuracy : 0.7500

'Positive' Class : 1

```
> #Ridge model w/lambda.1se
```

```
> confusionMatrix(as.factor(pred.ridge.1se))
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	72	27
1	18	63

Accuracy : 0.75  
 95% CI : (0.6801, 0.8114)  
 No Information Rate : 0.5  
 P-Value [Acc > NIR] : 6.046e-12

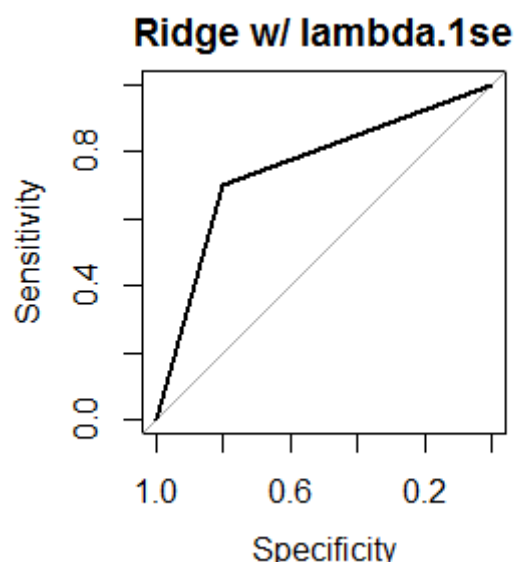
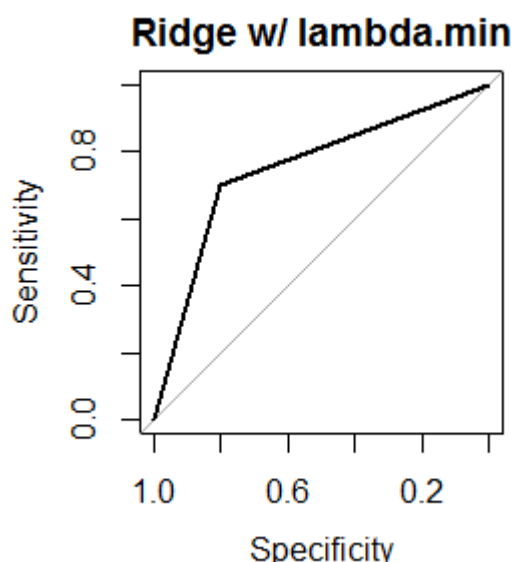
Kappa : 0.5

Mcnemar's Test P-value : 0.233

Sensitivity : 0.7000  
 Specificity : 0.8000  
 Pos Pred value : 0.7778  
 Neg Pred Value : 0.7273  
 Prevalence : 0.5000  
 Detection Rate : 0.3500  
 Detection Prevalence : 0.4500  
 Balanced Accuracy : 0.7500

'Positive' Class : 1

Although the two lambda parameters are different, they produced the same results. Both models score an overall accuracy of 0.75, and it can be noted again that the sensitivity is lower than the specificity. The ROC curves display similar imbalance, but a larger AUC:



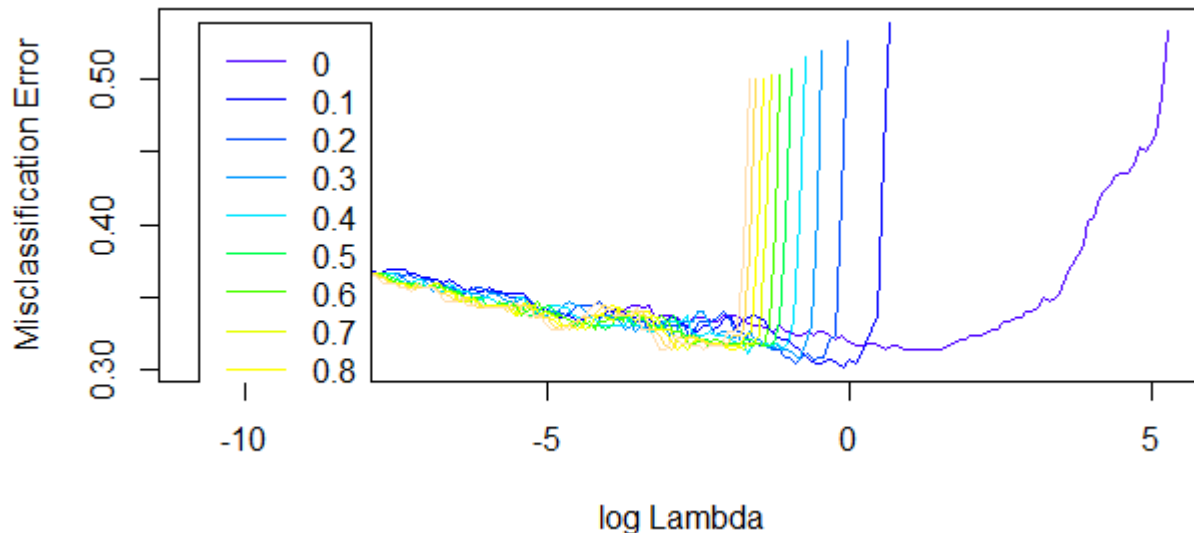
Therefore, the ridge models are more performant than the lasso models, at least from a predictive point of view. However, more testing should be performed to distinct between the two ridge models.

## Elastic Net:

To fit the elastic net model, both  $\alpha$  and  $\lambda$  have to be determined. First, the `cva.glmnet()` function has been used to find an optimal  $\alpha$  value, from 0 (ridge) to 1 (lasso):

```
> cv.net <- cva.glmnet(x, y, type.measure='class',  
+                      family='binomial', alpha=seq(0,1,0.1))
```

Plotting the glmnet object gives the following output:



According to this graph, it seems that  $\alpha=0.1$  is an optimum, as it has the lowest misclassification error rate around 0.30 when  $\log(\lambda)$  is around zero.

Thus, `cv.glmnet()` and  $\alpha=0.1$  have been used to fit the elastic net model. Here are the performance metrics and ROC curves that have been computed:

```
> #Elastic Net w/ lambda.min  
> confusionMatrix(as.factor(pred.net.min),  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	68	28
1	22	62

Accuracy : 0.7222  
95% CI : (0.6507, 0.7863)  
No Information Rate : 0.5  
P-Value [Acc > NIR] : 1.067e-09

Kappa : 0.4444

Mcnemar's Test P-Value : 0.4795

Sensitivity : 0.6889  
Specificity : 0.7556  
Pos Pred value : 0.7381  
Neg Pred value : 0.7083  
Prevalence : 0.5000  
Detection Rate : 0.3444  
Detection Prevalence : 0.4667  
Balanced Accuracy : 0.7222

'Positive' Class : 1

```
> #Elastic Net w/ lambda.1se  
> confusionMatrix(as.factor(pred.net.1se),  
Confusion Matrix and Statistics
```

	Reference	
Prediction	0	1
0	70	31
1	20	59

Accuracy : 0.7167  
95% CI : (0.6448, 0.7812)  
No Information Rate : 0.5  
P-Value [Acc > NIR] : 2.766e-09

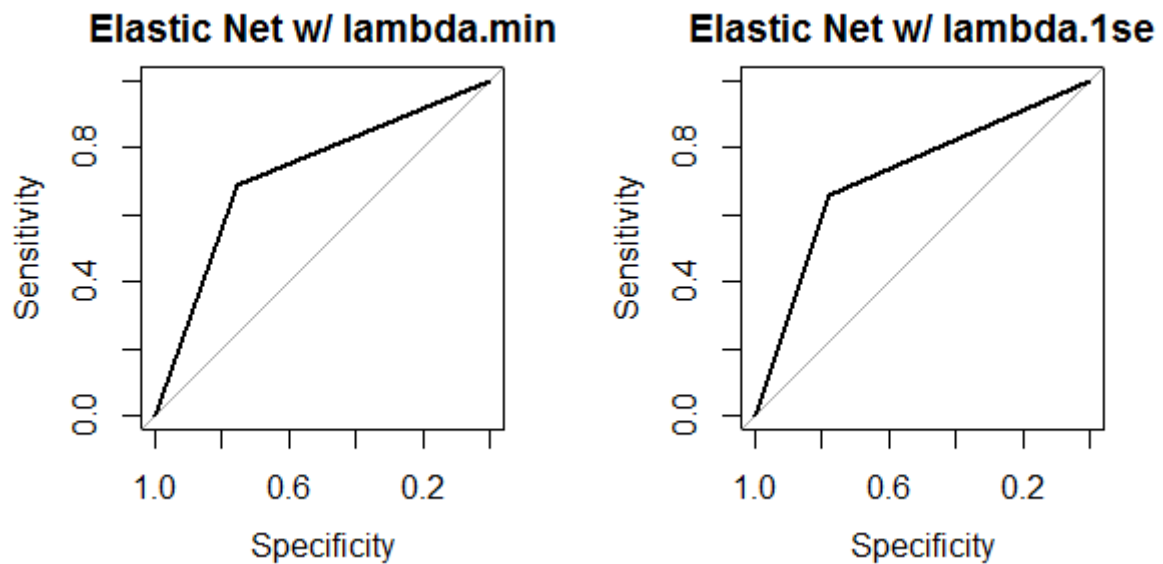
Kappa : 0.4333

Mcnemar's Test P-Value : 0.1614

Sensitivity : 0.6556  
Specificity : 0.7778  
Pos Pred value : 0.7468  
Neg Pred value : 0.6931  
Prevalence : 0.5000  
Detection Rate : 0.3278  
Detection Prevalence : 0.4389  
Balanced Accuracy : 0.7167

'Positive' Class : 1





#### 4. DISCUSSION

	Accuracy	Sensitivity	Specificity
Lasso w/ lambda.min	0.7222	0.6889	0.7556
Lasso w/ lambda.1se	0.7056	0.6444	0.7667
Ridge w/ lambda.min	0.75	0.7	0.8
Ridge w/ lambda.1se	0.75	0.7	0.8
Elastic Net w/ lambda.min	0.7222	0.6889	0.7556
Elastic Net w/ lambda.1se	0.7167	0.6556	0.7778

According to the results above, ridge regression (with *lambda.min* or *lambda.1se*) seems to be the most appropriate algorithm for predicting which NBA player will have a career length greater than 5 years – or not – with an overall accuracy of 0.75.

Several ameliorations could be brought to the model:

- Fine-tuning lambda between 0 and 0.1 since the elastic net model did not outperform the ridge one. The optimal lambda may lie between those two values.
- All of the models seem to underestimate the players' career length (low sensitivity). It could be interesting to find out why.
- Adding more observations to stabilise the classification rates and therefore obtaining more consistent models.